

## Core Concepts

### #### Monolithic vs Microservices

Monolithic: Entire application packaged into a single unit (difficult to scale, update).

Microservices: Application broken into smaller independent services (scalable, fault-tolerant).

Kubernetes is built to manage microservices at scale.

### #### Kubernetes Architecture

Control Plane: API Server, Scheduler, Controller Manager, ETCD (stores cluster state).

Worker Nodes: Run Pods, kubelet, kube-proxy, container runtime (Docker/Containerd).

Pods: Smallest deployable unit in Kubernetes.

### # Setup on Local/AWS EC2

Local: Using Minikube, Kind, or Docker Desktop.

Cloud: Deploying clusters on AWS EC2 (via Kops, EKS, kubeadm).

### # Kubectl

CLI tool for interacting with Kubernetes clusters. Example:

```
kubectl get pods
kubectl describe pod <pod-name>
kubectl apply -f deployment.yaml
```

### #### Pods

Smallest unit in Kubernetes (can hold one or more containers).

Each pod gets its own IP inside the cluster.

### #### Namespaces

Logical isolation within the cluster (e.g., dev, test, prod).

### #### Labels, Selectors, Annotations

Labels: Key-value pairs for grouping (e.g., app=frontend).

Selectors: Used to filter resources by labels.

Annotations: Metadata (not used for selection).

-----

### #### Workloads

These define how applications run in Kubernetes:

# Deployments – Manage stateless applications, handle scaling & rolling updates.

# StatefulSets – Manage stateful apps (databases, Kafka, etc.), provide stable

identity & storage.

# DaemonSets – Ensure a copy of a pod runs on all (or some) nodes (e.g., logging agents).

# ReplicaSets – Ensure a specified number of pod replicas are running.

# Jobs – Run tasks that complete and then exit (batch jobs).

# CronJobs – Scheduled jobs (like Linux cron, e.g., run backup every night).

-----  
-----

#### #### Networking

##### # Cluster Networking

Each Pod gets a unique IP.

Kubernetes ensures connectivity across nodes using CNI plugins (Calico, Flannel, Weave).

##### # Services

ClusterIP: Internal access only.

NodePort: Exposes service on each node's IP:port.

LoadBalancer: External traffic routing (uses cloud provider's LB).

Headless: No cluster IP, DNS directly maps to pod IPs.

##### # Ingress

Manages external access to services (HTTP/HTTPS).

Acts as a reverse proxy and load balancer with routing rules.

##### # Network Policies

Define which Pods can communicate with each other (like firewalls inside the cluster).

-----  
-----

#### #### Storage

##### # Persistent Volumes (PV)

Cluster-wide storage abstraction (can be backed by AWS EBS, NFS, GCP disk, etc.).

##### # Persistent Volume Claims (PVC)

Pod request for storage (claims PV).

##### # StorageClasses

Define how dynamic storage should be provisioned (e.g., SSD vs HDD).

##### # ConfigMaps

Store configuration data (non-sensitive). Example: environment variables.

# Secrets

Store sensitive information (passwords, API keys, TLS certs).

Base64 encoded, not encrypted by default.

-----

---

# 🧑💻 Kubernetes Interview Q\&A (Advanced Topics)

---

## \*\*1. Scaling and Scheduling\*\*

### ? What is the difference between HPA and VPA?

**Answer:**

\* **HPA (Horizontal Pod Autoscaler)** scales the number of pod replicas based on CPU, memory, or custom metrics.

\* **VPA (Vertical Pod Autoscaler)** adjusts the **resources (CPU/memory)** of existing pods.

👉 HPA = scale **out/in**, VPA = scale **up/down**.

**Example:**

```bash

kubectl autoscale deployment nginx --cpu-percent=50 --min=2 --max=10

```

---

### ? How do Node Affinity and Taints/Tolerations differ?

**Answer:**

\* **Node Affinity**: Soft/hard rules to **attract pods** to specific nodes (e.g., "run this on SSD nodes").

\* **Taints/Tolerations**: Mechanism to **repel pods** unless they tolerate the taint.

👉 Affinity = "go there", Taint = "keep away unless tolerated".

---

### ? What are Resource Quotas and Limits?

**Answer:**

\* **Resource Requests**: Minimum CPU/memory a pod needs.

\* **Limits**: Max resources a pod can consume.

\* **ResourceQuota**: Enforce per-namespace limits.

**Example:** Prevents one team from hogging cluster resources.

---

### ? What are Probes in Kubernetes?

**\*\*Answer:\*\***

- \* **\*\*Liveness Probe\*\***: Checks if the container is alive; restarts if unhealthy.
- \* **\*\*Readiness Probe\*\***: Checks if container is ready to serve traffic.
- \* **\*\*Startup Probe\*\***: Used for slow-starting apps.

---

**## \*\*2. Cluster Administration\*\***

**### ? Explain RBAC in Kubernetes.**

**\*\*Answer:\*\***

RBAC (Role-Based Access Control) restricts what users or service accounts can do.

- \* **\*\*Role\*\***: Namespace-scoped permissions.
- \* **\*\*ClusterRole\*\***: Cluster-wide permissions.
- \* **\*\*RoleBinding/ClusterRoleBinding\*\***: Bind roles to users or service accounts.

👉 Example: Only allow a dev to `get` and `list` pods in `dev` namespace.

---

**### ? What is a CRD (Custom Resource Definition)?**

**\*\*Answer:\*\***

- \* CRDs allow extending Kubernetes with custom objects.
- \* Example: Create a custom resource `Database` to manage MySQL clusters like any native Kubernetes object.

👉 This is the foundation for building **\*\*Operators\*\***.

---

**### ? How do you upgrade a Kubernetes cluster?**

**\*\*Answer:\*\***

- \* In **\*\*managed services (EKS/GKE/AKS)\*\***: Use provider's upgrade command/console.
- \* In **\*\*self-managed cluster\*\***:

1. Upgrade `kubeadm`
2. Upgrade control-plane nodes (`kubeadm upgrade`)
3. Upgrade worker nodes (`kubectl drain`, update kubelet, uncordon).

---

**## \*\*3. Monitoring and Logging\*\***

**### ? How do you monitor resource usage in Kubernetes?**

**\*\*Answer:\*\***

- \* Install **\*\*Metrics Server\*\*** → enables `kubectl top nodes/pods`.
- \* For detailed monitoring, use **\*\*Prometheus + Grafana\*\***.

---

**### ? How do you handle logging in Kubernetes?**

**\*\*Answer:\*\***

- \* Default: `kubectl logs <pod>`.

\* For centralized logging: use **EFK** (Elasticsearch + Fluentd + Kibana) or **Loki + Promtail + Grafana**.

---

## ## **4. Advanced Features**

### ? What are Operators in Kubernetes?

**Answer:**

\* Operators extend Kubernetes to manage complex apps (e.g., databases).  
\* They use CRDs + controllers to automate lifecycle tasks like backup, failover, scaling.

---

### ? What is Helm and why is it used?

**Answer:**

\* Helm is the **package manager for Kubernetes**.  
\* It bundles multiple YAML manifests into a reusable **chart**.  
👉 Instead of applying 10 different YAMLS, you deploy 1 Helm chart.

---

### ? What is a Service Mesh?

**Answer:**

\* A Service Mesh (Istio, Linkerd) manages **service-to-service communication**.  
\* Features: observability, traffic routing, retries, security (mTLS).  
👉 Useful for microservices.

---

## ## **5. Security**

### ? What are Pod Security Standards (PSS)?

**Answer:**

\* Defines security levels for pods:  
  \* **Privileged** → full access (least secure).  
  \* **Baseline** → minimal restrictions.  
  \* **Restricted** → strong security (no root user, limited privileges).

---

### ? How do you secure container images in Kubernetes?

**Answer:**

\* Use **image scanning tools** (Trivy, Anchore).  
\* Ensure images are signed (Cosign, Notary).  
\* Use private registries with authentication.

---

### ? How do you secure secrets in Kubernetes?

**Answer:**

- \* Store sensitive data in **Secrets** (base64 encoded).
- \* Enable **encryption at rest** using KMS providers (AWS KMS, GCP KMS).
- \* Mount secrets as env vars or files.

---

## ## **6. Cloud-Native Kubernetes**

### ? What's the difference between EKS, AKS, and GKE?

**Answer:**

- \* All are **managed Kubernetes services**.
- \* **EKS (AWS)**, **AKS (Azure)**, **GKE (Google Cloud)** → manage the control plane for you.
- \* You only manage worker nodes & workloads.

---

### ? How does Cluster Autoscaler work?

**Answer:**

- \* Scales worker nodes **up/down** depending on pending pods.
- \* Works with cloud provider APIs (EC2 Auto Scaling Groups, GCP Instance Groups).

---

### ? What are Spot/Preemptible Nodes?

**Answer:**

- \* **Spot (AWS)** / **Preemptible (GCP)** are cheap, short-lived nodes.
- \* Great for **batch jobs** or **non-critical workloads**.
- \* If the cloud provider needs capacity, your nodes may be terminated.

---

-----

Great question, Suraj 🙌 – Kubernetes **debugging & troubleshooting** is one of the **most asked interview topics**, and also the most important in real-world clusters. I'll cover it in a structured way with **common issues, commands, YAML fixes, and interview-ready answers**.

---

## # 🛠️ Kubernetes Debugging & Troubleshooting

---

### ## **1. Pod Issues**

### ? My Pod is stuck in `Pending`. How do you debug?

**Answer:**

- \* Check if there are enough resources (CPU/Memory) on nodes.
- \* Verify **PVC binding** if using persistent storage.
- \* Look at events.

**Commands:**

```
```bash
kubectl get pods
kubectl describe pod <pod-name>
kubectl get events --sort-by=.metadata.creationTimestamp
```
```

**\*\*Common Fix:\*\***

- \* If node lacks resources → scale cluster.
- \* If PVC not bound → check PV and StorageClass.

---

### ? My Pod is in `CrashLoopBackOff`. How do you debug?

**\*\*Answer:\*\***

- \* Container keeps failing on startup.
- \* Check logs of the container.

**\*\*Commands:\*\***

```
```bash
kubectl logs <pod-name>
kubectl logs <pod-name> -c <container-name> # multi-container pod
kubectl describe pod <pod-name>
```
```

**\*\*Fixes:\*\***

- \* Wrong image? → Fix `image` field.
- \* Misconfigured environment variables? → Check ConfigMaps/Secrets.
- \* App needs time? → Add **\*\*readiness/liveness probes\*\*** carefully.

---

### ? How do you get inside a running Pod for debugging?

```
```bash
kubectl exec -it <pod-name> -- /bin/sh
kubectl exec -it <pod-name> -- /bin/bash
```
```

👉 Useful to check app logs/configs inside the container.

---

## **\*\*2. Deployment & Replica Issues\*\***

### ? My Deployment is not scaling. How do you debug?

- \* Check replica count:

```
```bash
kubectl get deploy
kubectl describe deploy <name>
```
```

- \* Check if HPA is working:

```
```bash
kubectl get hpa
```
```

**\*\*Common Fixes:\*\***

- \* Metrics server not installed → HPA won't work.
- \* Node resource shortage.

---

### ? How do you roll back a failed Deployment?

```
```bash
kubectl rollout undo deployment <deployment-name>
```
```

---

## **\*\*3. Service & Networking Issues\*\***

### ? My Service is not reachable. How do you debug?

1. Check Service:

```
```bash
kubectl get svc
kubectl describe svc <svc-name>
```
```

2. Check Endpoints (should map to pods):

```
```bash
kubectl get endpoints <svc-name>
```
```

3. Exec into a pod and curl the service name:

```
```bash
kubectl exec -it <pod> -- curl http://<svc-name>:<port>
```
```

**\*\*Common Fixes:\*\***

- \* Service selector labels don't match Pod labels.
- \* Pod not running on correct port.

---

### ? My Ingress is not working. How do you debug?

1. Check Ingress rules:

```
```bash
kubectl describe ingress <name>
```
```

2. Verify Ingress Controller is deployed (`nginx-ingress`, `traefik`, etc.).
3. DNS must resolve to Ingress external IP.

---

## **\*\*4. Node Issues\*\***

### ? A Node is in `NotReady` state. What do you check?

\* Node status:

```
```bash
kubectl get nodes
kubectl describe node <node-name>
```
```



```

` ``
* Kubelet logs on the node.
* Check `docker` or `containerd` runtime.
* Ensure network (CNI plugin) is running.

**Fix:** Restart kubelet, check disk/memory, check CNI pods.

```

---

## ## \*\*5. Storage Issues\*\*

### ? My Pod is stuck waiting for PVC. How do you debug?

```

` ``bash
kubectl get pvc
kubectl describe pvc <pvc-name>
kubectl get pv
` ``

```

\*\*Fix:\*\*

\* Ensure a matching PV exists or dynamic provisioning via StorageClass works.

---

## ## \*\*6. RBAC Issues\*\*

### ? I get "forbidden" errors while using kubectl. Why?

\*\*Answer:\*\*

\* RBAC rules don't allow your user/service account to perform that action.

\*\*Debug:\*\*

```

` ``bash
kubectl auth can-i get pods --as=developer
` ``

```

\*\*Fix:\*\* Add appropriate `Role`/`RoleBinding`.

---

## ## \*\*7. Debugging Tools & Commands\*\*

\* \*\*Check pod logs\*\*

```

` ``bash
kubectl logs <pod>
` ``

```

\* \*\*Check events in namespace\*\*

```

` ``bash
kubectl get events -n <namespace>
` ``

```

\* \*\*Debug pods without a shell\*\*

```

` ``bash
kubectl run tmp --rm -it --image=busybox -- sh
` ``

```

👉 Create a temporary debug pod.

\* **\*\*Describe resources\*\***

```
```bash
kubectl describe pod <pod>
kubectl describe node <node>
```
```

\* **\*\*Network testing pod\*\***


```
```bash
kubectl run net-test --image=busybox:1.28 -it --rm -- nslookup <service-name>
```
```

---

## **\*\*8. Interview-Worthy Troubleshooting Scenarios\*\***

1. **\*\*Pod keeps restarting\*\*** → Misconfigured app/env variables.
2. **\*\*Service has no endpoints\*\*** → Labels mismatch.
3. **\*\*Ingress not routing\*\*** → Missing Ingress Controller.
4. **\*\*PVC not bound\*\*** → No suitable PV/StorageClass.
5. **\*\*Forbidden errors\*\*** → RBAC issue.
6. **\*\*Node NotReady\*\*** → Kubelet/CNI failure.

---

 With these debugging flows, you'll show interviewers that you can **\*\*systematically troubleshoot\*\*** Kubernetes like a real SRE/DevOps engineer.

Would you like me to now create a **\*\*flowchart-style checklist (step-by-step debugging guide)\*\*** so you can quickly recall during interviews or on the job?